

# 시계열 분석

## 라이브러리 호출

```
In [68]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib
import matplotlib.pyplot as plt
plt.style.use('seaborn-whitegrid')
%matplotlib inline
pd.options.display.max_columns = None

import os

from sklearn.model_selection import train_test_split
from sklearn.metrics import r2_score

import statsmodels.api as sm
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.arima_model import ARIMA
from statsmodels.tsa.statespace.sarimax import SARIMAX
from pmdarima.arima import auto_arima  ## ADP 볼 때는 없을 패키지

import itertools # 내장 패키지
```

## 그래프 한글 깨짐 방지

```
In [69]: from matplotlib import font_manager, rc
path = 'malgun.ttf'
font_name = font_manager.FontProperties(fname=path).get_name()
rc('font', family=font_name)
```

## 데이터 로딩

```
In [70]: df = pd.read_csv('./data/bikeshare.csv')
```

## 데이터 구조 확인

In [71]: `df.head()`

Out[71]:

	datetime	season	holiday	workingday	weather	temp	atemp	humidity	windspeed	casual
0	2011-01-01 0:00	A	0	0	1	9.84	14.395	81	0.0	3
1	2011-01-01 1:00	A	0	0	1	9.02	13.635	80	0.0	8
2	2011-01-01 2:00	A	0	0	1	9.02	13.635	80	0.0	5
3	2011-01-01 3:00	A	0	0	1	9.84	14.395	75	0.0	3
4	2011-01-01 4:00	A	0	0	1	9.84	14.395	75	0.0	0

In [72]: `df.tail()`

Out[72]:

	datetime	season	holiday	workingday	weather	temp	atemp	humidity	windspeed	casual
10881	2012-12-19 19:00	D	0	1	1	15.58	19.695	50	26.0027	
10882	2012-12-19 20:00	D	0	1	1	14.76	17.425	57	15.0013	
10883	2012-12-19 21:00	D	0	1	1	13.94	15.910	61	15.0013	
10884	2012-12-19 22:00	D	0	1	1	13.94	17.425	61	6.0032	
10885	2012-12-19 23:00	D	0	1	1	13.12	16.665	66	8.9981	

In [73]: `df.shape`

Out[73]: (10886, 12)

In [74]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10886 entries, 0 to 10885
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   datetime        10886 non-null  object
1   season          10886 non-null  object
2   holiday         10886 non-null  int64
3   workingday      10886 non-null  int64
4   weather         10886 non-null  int64
5   temp            10886 non-null  float64
6   atemp           10886 non-null  float64
7   humidity        10886 non-null  int64
8   windspeed       10886 non-null  float64
9   casual          10886 non-null  int64
10  registered      10886 non-null  int64
11  count           10886 non-null  int64
dtypes: float64(3), int64(7), object(2)
memory usage: 1020.7+ KB
```

## 날짜 데이터 전처리

### 날짜 데이터를 Timestamp 형식으로 변환

```
In [75]: # 날짜 형식이 연, 월, 일, 시, 분, 초 형태일 때
df['datetime'] = pd.to_datetime(df['datetime'])

# pd.to_datetime('2012-12-19 20:00', format='%Y-%m-%d %H:%M')
```

===== 참고 : 날짜형식이 13자리 숫자일 때 대비

```
In [76]: # 13자리 숫자일 때
import datetime
timestamp = 1463460958000
datetimeobj = datetime.datetime.fromtimestamp(timestamp/1000)
print(datetimeobj, type(datetimeobj))
# 이후에 pd.to_datetime으로 변환
a = pd.to_datetime(datetimeobj)
print(a, type(a))

2016-05-17 13:55:58 <class 'datetime.datetime'>
2016-05-17 13:55:58 <class 'pandas._libs.tslibs.timestamps.Timestamp'>
```

```
In [77]: # datetime to timestamp
import time
timestamp = time.mktime(datetimeobj.timetuple())
timestamp
```

```
Out[77]: 1463460958.0
```

```
In [78]: import datetime
datetime.date(year=2019, month=10, day=1)
```

```
Out[78]: datetime.date(2019, 10, 1)
```

===== 여기까지

## 날짜 데이터로부터 연, 월, 일, 시, 요일 데이터 추출

```
In [79]: df['year'] = df['datetime'].map(lambda x: x.year) # 연
df['month'] = df['datetime'].map(lambda x: x.month) # 월
df['day'] = df['datetime'].map(lambda x: x.day) # 일
df['hour'] = df['datetime'].map(lambda x: x.hour) # 시
df['dayofweek'] = df['datetime'].map(lambda x: x.dayofweek) # 요일
```

## 날짜 데이터를 시 기준으로 그룹핑( 다른 데이터는 평균값 계산 ) 후 인덱스로 설정

```
In [80]: # 그룹핑
df = df.groupby(['year', 'month', 'day', 'hour'])['temp', 'humidity', 'windspeed',
'count'].mean().reset_index()
```

C:\Users\50008313\AppData\Local\Continuum\anaconda3\lib\site-packages\ipykernel\_launcher.py:2: FutureWarning: Indexing with multiple keys (implicitly converted to a tuple of keys) will be deprecated, use a list instead.

```
In [81]: # 다시 date 컬럼 만들어주기
df['date'] = df['year'].astype('str') + '-' + df['month'].astype('str') + '-' + df['day'].astype('str') + '-' + df['hour'].astype('str')

# date 컬럼 형식 변경 -> datetime
df['date'] = pd.to_datetime(df['date'], format='%Y-%m-%d-%H')
```

```
In [82]: df = df.set_index('date')
```

In [83]: df.head(3)

Out[83]:

	year	month	day	hour	temp	humidity	windspeed	count
date								
2011-01-01 00:00:00	2011	1	1	0	9.84	81	0.0	16
2011-01-01 01:00:00	2011	1	1	1	9.02	80	0.0	40
2011-01-01 02:00:00	2011	1	1	2	9.02	80	0.0	32

In [84]: # 필요 컬럼만 선택  
df = df[['temp', 'humidity', 'windspeed', 'count']]

In [85]: df.head()

Out[85]:

	temp	humidity	windspeed	count
date				
2011-01-01 00:00:00	9.84	81	0.0	16
2011-01-01 01:00:00	9.02	80	0.0	40
2011-01-01 02:00:00	9.02	80	0.0	32
2011-01-01 03:00:00	9.84	75	0.0	13
2011-01-01 04:00:00	9.84	75	0.0	1

In [116]: # 시계열용 데이터는 따로 빼둬  
dfts = pd.DataFrame(df['count'])

In [117]: dfts

Out[117]:

date	
2011-01-01 00:00:00	16.0
2011-01-01 01:00:00	40.0
2011-01-01 02:00:00	32.0
2011-01-01 03:00:00	13.0
2011-01-01 04:00:00	1.0
...	
2012-12-19 19:00:00	336.0
2012-12-19 20:00:00	241.0
2012-12-19 21:00:00	168.0
2012-12-19 22:00:00	129.0
2012-12-19 23:00:00	88.0

Name: count, Length: 10886, dtype: float64

===== 여기부터는 y값 외에독립변수가 더 있을 경우

===== 시계열 모델만 만들 거면 비시계열 모델링 이후로 이동

## 데이터 타입 맞춰주기 =====

```
In [86]: df.columns
```

```
Out[86]: Index(['temp', 'humidity', 'windspeed', 'count'], dtype='object')
```

```
In [87]: col_id = []  
col_cat = []  
col_int = []  
col_float = ['temp', 'humidity', 'windspeed', 'count']  
col_bool = []  
col_num = col_int + col_float
```

```
In [88]: df[col_cat] = df[col_cat].astype('str')  
df[col_int] = df[col_int].astype('int', errors = 'ignore')  
df[col_float] = df[col_float].astype('float')
```

## DQ Check(빈도분석, 분포분석)

### 연속형 변수

```

In [89]: def DA(data):
    da = data.describe(percentiles=[0.05, 0.1, 0.25, 0.5, 0.75, 0.9, 0.95])
    da = da.T
    df1 = data.isna().sum()
    df1.name = 'missing'
    df2 = data.median()
    df2.name = 'median'
    df3 = np.var(data)
    df3.name = 'variance'
    df4 = data.skew()
    df4.name = 'skewness'
    df5 = data.kurtosis()
    df5.name = 'kurtosis'

    da = pd.concat([da, df1, df2, df3, df4, df5], axis =1)
    da['total'] = da['count'] + da['missing']

    col_nm = da.columns.tolist()
    order = ['total', 'count', 'missing', 'mean', 'median', 'std', 'variance', 'skewness', 'kurtosis', 'min',
              '5%', '10%', '25%', '50%', '75%', '90%', '95%', 'max']
    col_nm_new=[]
    for i in order:
        col_nm_new.append(i)
    da = da[col_nm_new]
    da = da.round(2)
    return da

```

```

In [90]: DA1 = DA(df[col_num])
DA1

```

Out[90]:

	total	count	missing	mean	median	std	variance	skewness	kurtosis	n
temp	10886.0	10886.0	0	20.23	20.5	7.79	60.70	0.00	-0.91	0.
humidity	10886.0	10886.0	0	61.89	62.0	19.25	370.34	-0.09	-0.76	0.
windspeed	10886.0	10886.0	0	12.80	13.0	8.16	66.65	0.59	0.63	0.
count	10886.0	10886.0	0	191.57	145.0	181.14	32810.30	1.24	1.30	1.

## 범주형 변수

```
In [91]: def DA_cat(data, col_cat):
        DA_cat = pd.DataFrame()

        for i in col_cat:
            a = data[i].value_counts(dropna=False).to_frame().sort_index().rename(
                columns={i:'count'}).reset_index()
            a['col_nm'] = i
            a = a.rename(columns = {'index':'class'})
            a = a[['col_nm', 'class', 'count']]
            b=data[i].value_counts(dropna = False, normalize = True).to_frame().so
            rt_index().rename(
                columns = {i:'ratio'}).reset_index()
            b = b['ratio'].to_frame()
            b['ratio'] = b['ratio'].round(2)
            c = pd.concat([a,b], axis = 1)
            DA_cat = pd.concat([DA_cat, c], axis=0)
        DA_cat = DA_cat.reset_index(drop=True)
        return DA_cat
```

```
In [92]: DA2 = DA_cat(df,col_cat+col_bool)
        DA2
```

```
Out[92]:
```

—

## 전처리(중복값, 결측치, 이상치 처리)

### 중복값



```
In [93]: df[df.duplicated(keep=False)].sort_values(['temp', 'humidity', 'windspeed', 'count'])
```

Out[93]:

	temp	humidity	windspeed	count
date				
2012-01-04 02:00:00	0.82	34.0	19.0012	1.0
2012-01-04 03:00:00	0.82	34.0	19.0012	1.0
2011-01-09 04:00:00	3.28	53.0	12.9980	1.0
2011-01-09 05:00:00	3.28	53.0	12.9980	1.0
2011-02-10 05:00:00	4.92	50.0	15.0013	6.0
...	...	...	...	...
2012-08-08 03:00:00	28.70	84.0	0.0000	7.0
2012-08-08 04:00:00	28.70	84.0	0.0000	7.0
2011-08-07 05:00:00	28.70	89.0	12.9980	5.0
2012-07-09 02:00:00	28.70	89.0	12.9980	5.0
2012-07-09 03:00:00	28.70	89.0	12.9980	5.0

161 rows × 4 columns

```
In [ ]: df.drop_duplicates()
df.drop_duplicates(['col1'], keep='last')
```

## 결측치

```
In [94]: df.isna().sum()
```

```
Out[94]: temp      0
humidity    0
windspeed   0
count       0
dtype: int64
```

```

In [ ]: # na 처리 : dropna(), fillna()
df.dropna() # nan이 하나라도 들어간 행은 삭제
df.dropna(how = 'all') # 데이터가 모두 nan인 행만 삭제 / 초기값: 'any'
## Parameters
# axis = 'index' / 'columns'
# subset = ['col1', 'col2', ...] # 적용 대상 컬럼 특정

df.fillna(0) # na를 0으로 채우기

new_data = {'a':0, 'b':1, 'c':-999}
df.fillna(new_data) # na 발생 시 a 열에는 0, b 열에는 1, c 열에는 -999로 채움
df.fillna(new_data, limit = 2) # 각 열별로 2개의 nan까지 대체
df.fillna(method = 'ffill') # 열 별로 바로 앞의 데이터로 채움
df.fillna(method = 'bfill') # 열 별로 바로 뒤의 데이터로 채움
# ffill의 경우 첫 행이거나, 앞의 데이터가 nan일 경우 nan유지. bfill도 반대로 동일

# 평균값, 중앙값으로 대체
df.loc[19, 'Leaflets'] = df['Leaflets'].mean() # 평균값으로
df.loc[19, 'Leaflets'] = df['Leaflets'].median() # 중앙값으로

```

## 이상치

```

In [95]: tmp = 'windspeed'

```

```


In [96]: sns.boxplot(y = tmp, data = df, orient = 'h')

```

```

Out[96]: <AxesSubplot:xlabel='windspeed'>

```



```

In [97]: # IQR 활용
q1 = df[tmp].quantile(.25)
q3 = df[tmp].quantile(.75)
iqr = q3-q1
min_iqr = q1 - 1.5 * iqr
max_iqr = q3 + 1.5 * iqr
min_from_all = df[tmp].min()
max_from_all = df[tmp].max()
if (min_iqr < min_from_all) :
    min_iqr = min_from_all
if (max_iqr > max_from_all) :
    max_iqr = max_from_all

outlier = df[(df[tmp] < min_iqr ) | (df[tmp] > max_iqr)] # 이상치 조회
outlier_index = outlier.index
print(outlier.shape)
outlier

```

(227, 4)

Out[97]:

	temp	humidity	windspeed	count
date				
2011-01-08 14:00:00	8.20	32.0	32.9975	95.0
2011-01-08 17:00:00	6.56	37.0	36.9974	69.0
2011-01-09 09:00:00	4.92	46.0	35.0008	19.0
2011-01-09 11:00:00	6.56	40.0	35.0008	49.0
2011-01-12 12:00:00	8.20	47.0	39.0007	55.0
...	...	...	...	...
2012-11-02 14:00:00	16.40	40.0	32.9975	262.0
2012-11-08 12:00:00	16.40	24.0	32.9975	235.0
2012-11-13 01:00:00	18.04	88.0	43.0006	5.0
2012-12-05 14:00:00	19.68	33.0	32.9975	218.0
2012-12-18 15:00:00	18.86	44.0	32.9975	246.0

227 rows × 4 columns

**min/max값으로 보정**

```

In [ ]: df.loc[(df[tmp] < min_iqr ),tmp] = min_iqr # 이상치 보정 - 하한치로 보정
df.loc[(df[tmp] > max_iqr ),tmp] = max_iqr # 이상치 보정 - 상한치로 보정

```

**이상치 제거**

```
In [ ]: df = df.drop(outlier_index, axis=0)
df.shape
```

## 파생변수 생성

```
In [98]: today = pd.to_datetime('2020-12-13')
```

```
In [99]: # Recency
cond1 = (today-df.index) >= pd.Timedelta('3000 days')
cond2 = ((today-df.index) < pd.Timedelta('3000 days')) & ((today-df.index) >= pd
.Timedelta('2000 days'))
cond3 = (today-df.index) < pd.Timedelta('2000 days')

df.loc[cond1, 'Recency'] = 1
df.loc[cond2, 'Recency'] = 2
df.loc[cond3, 'Recency'] = 3
```

```
In [100]: # Frequency
df.loc[df['count'] <= 10, 'Frequency'] = 1
df.loc[(df['count'] > 10) & (df['count'] <= 20), 'Frequency'] = 2
df.loc[df['count'] > 20, 'Frequency'] = 3
```

```
In [101]: # Monetary
df['Monetary'] = df['count'] * df['temp']
```

```
In [102]: df.head(3)
```

Out[102]:

	temp	humidity	windspeed	count	Recency	Frequency	Monetary
date							
2011-01-01 00:00:00	9.84	81.0	0.0	16.0	1.0	2.0	157.44
2011-01-01 01:00:00	9.02	80.0	0.0	40.0	1.0	3.0	360.80
2011-01-01 02:00:00	9.02	80.0	0.0	32.0	1.0	3.0	288.64

## 데이터 마트 DQ Check, 변수선택 및 EDA

### DQ Check

```
In [103]: df.columns
```

```
Out[103]: Index(['temp', 'humidity', 'windspeed', 'count', 'Recency', 'Frequency',
'Monetary'],
dtype='object')
```

```
In [104]: col_num = ['temp', 'humidity', 'windspeed', 'count', 'Monetary']
col_cat = ['Recency', 'Frequency']
```

```
In [105]: DA3 = DA(df[col_num])
DA3
```

Out[105]:

	total	count	missing	mean	median	std	variance	skewness	kurtosis
<b>temp</b>	10886.0	10886.0	0	20.23	20.50	7.79	60.70	0.00	-0.
<b>humidity</b>	10886.0	10886.0	0	61.89	62.00	19.25	370.34	-0.09	-0.
<b>windspeed</b>	10886.0	10886.0	0	12.80	13.00	8.16	66.65	0.59	0.
<b>count</b>	10886.0	10886.0	0	191.57	145.00	181.14	32810.30	1.24	1.
<b>Monetary</b>	10886.0	10886.0	0	4432.39	2629.74	5023.07	25228922.33	1.66	2.

```
In [106]: DA4 = DA_cat(df, col_cat)
DA4
```

Out[106]:

	col_nm	class	count	ratio
<b>0</b>	Recency	1.0	9519	0.87
<b>1</b>	Recency	2.0	1367	0.13
<b>2</b>	Frequency	1.0	1229	0.11
<b>3</b>	Frequency	2.0	631	0.06
<b>4</b>	Frequency	3.0	9026	0.83

## 변수 제외

```
In [107]: df = df.drop(columns = ['Recency'], axis=1)
```


## EDA

```
In [ ]: # 범주형 x별 y의 평균
sns.barplot(x='season', y='windspeed', data=df)
```

```
In [ ]: # 범주형(또는 가지수가 많지 않은 연속형) 변수의 데이터별 count
sns.countplot(y='holiday', data=df)
```


```
In [108]: sns.lineplot(x = df.index, y = 'temp', data = df)
```

```
Out[108]: <AxesSubplot:xlabel='date', ylabel='temp'>
```



```
In [109]: sns.scatterplot(x = 'temp', y = 'count', data = df)
```


```
Out[109]: <AxesSubplot:xlabel='temp', ylabel='count'>
```



## 종속변수 분포 확인 및 전처리

```
In [110]: df['count'].plot(kind='hist')
```


```
Out[110]: <AxesSubplot:ylabel='Frequency'>
```



```
In [111]: df['y2'] = np.log1p(df['count']) # inverse 는 np.expm1()
```

```
In [112]: df['y2'].plot(kind='hist')
```

```
Out[112]: <AxesSubplot:ylabel='Frequency'>
```



In [113]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 10886 entries, 2011-01-01 00:00:00 to 2012-12-19 23:00:00
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   temp        10886 non-null   float64
1   humidity     10886 non-null   float64
2   windspeed    10886 non-null   float64
3   count        10886 non-null   float64
4   Frequency    10886 non-null   float64
5   Monetary     10886 non-null   float64
6   y2           10886 non-null   float64
dtypes: float64(7)
memory usage: 680.4 KB
```

In [114]: df.head(3)

Out[114]:

	temp	humidity	windspeed	count	Frequency	Monetary	y2
date							
2011-01-01 00:00:00	9.84	81.0	0.0	16.0	2.0	157.44	2.833213
2011-01-01 01:00:00	9.02	80.0	0.0	40.0	3.0	360.80	3.713572
2011-01-01 02:00:00	9.02	80.0	0.0	32.0	3.0	288.64	3.496508

```
In [118]: import statsmodels.api as sm
from patsy import dmatrices
```

```
y, X = dmatrices('y2 ~ temp + humidity + windspeed + Frequency + Monetary', data=df, return_type='dataframe')
```

## VIF 확인 필요 (y값 섞여들어가지 않게 주의!!)



```
In [119]: from statsmodels.stats.outliers_influence import variance_inflation_factor

vif = pd.DataFrame()
vif['VIF Factor'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
vif['features'] = X.columns
vif
```

Out[119]:

	VIF Factor	features
0	47.765907	Intercept
1	1.576327	temp
2	1.257986	humidity
3	1.116640	windspeed
4	1.180948	Frequency
5	1.860919	Monetary

## train, test split

```
In [120]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=0)
```

## StandardScaler

```
In [121]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
```

```
In [122]: scaler.fit(X_train)
X_train_scale = scaler.transform(X_train)
X_test_scale = scaler.transform(X_test)
```

```
In [123]: # 컬럼명 다시 붙여주기
X_train_scale = pd.DataFrame(X_train_scale, columns = X_train.columns)
X_test_scale = pd.DataFrame(X_test_scale, columns = X_test.columns)
```

## 군집화 수행


```
In [124]: # X_train_scale, X_test_scale, y_train, y_test가 현재 변수

from sklearn.cluster import KMeans

def elbow(X):
    sse = []
    for i in range(1, 11) :
        km = KMeans(n_clusters=i, init='k-means++', random_state = 0)
        km.fit(X)
        sse.append(km.inertia_)

    plt.plot(range(1, 11), sse, marker='o')
    plt.xlabel('n_clusters')
    plt.ylabel('SSE')
    plt.show()
```

```
In [125]: elbow(X_train_scale)
```



```
In [126]: from sklearn.metrics import silhouette_samples, silhouette_score

def sil(X):
    si = [] # 실루엣계수
    for i in range(2,11): # cluster가 2개인것 부터 10개까지!!!!
        km = KMeans(n_clusters=i, init='k-means++', random_state=0)
        km.fit(X)
        si.append(silhouette_score(X, km.labels_))
    print(np.round(si,3))
sil(X_train_scale)
```

```
[0.24  0.282 0.283 0.248 0.236 0.229 0.22  0.219 0.213]
```

## 군집 수 직접 지정해서 군집화

```
In [127]: kmeans = KMeans(n_clusters=4, init='k-means++', max_iter=300, random_state=0)
kmeans.fit(X_train_scale)
```

```
Out[127]: KMeans(n_clusters=4, random_state=0)
```


## 군집화 결과 프로파일링

```
In [128]: # 스케일링 풀고 프로파일링

df_profile = pd.DataFrame(scaler.inverse_transform(X_train_scale), columns = X_train.columns)
df_profile['kmeans'] = kmeans.labels_
```


```
In [130]: sns.barplot(df_profile['kmeans'], df_profile['windspeed'])
```

```
Out[130]: <AxesSubplot:xlabel='kmeans', ylabel='windspeed'>
```



```
In [132]: sns.barplot(df_profile['kmeans'], df_profile['Monetary'])
```

```
Out[132]: <AxesSubplot:xlabel='kmeans', ylabel='Monetary'>
```



## 구진화 결과를 새로운 컬럼으로 추가(train, test 모두 수행)

```
In [133]: X_train_scale['kmeans'] = kmeans.labels_
```

```
In [134]: kmeans_test = kmeans.predict(X_test_scale)
X_test_scale['kmeans'] = kmeans_test
```

## 모델링

```
In [136]: from sklearn.linear_model import Ridge, Lasso, ElasticNet, HuberRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.neighbors import KNeighborsRegressor
from sklearn.kernel_ridge import KernelRidge
from sklearn.neural_network import MLPRegressor
from sklearn.svm import SVR

from sklearn.model_selection import GridSearchCV, train_test_split, KFold, cross_val_score
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
```

```
In [137]: models = []
models.append(('Ridge', Ridge()))
models.append(('Lasso', Lasso()))
models.append(('ElasticNet', ElasticNet()))
models.append(('Huber', HuberRegressor()))
models.append(('DT', DecisionTreeRegressor()))
models.append(('RF', RandomForestRegressor()))
models.append(('KNN', KNeighborsRegressor()))
models.append(('KernelRidge', KernelRidge()))
models.append(('MLP', MLPRegressor()))
models.append(('SVR', SVR()))
```

```
In [138]: models
```

```
Out[138]: [('Ridge', Ridge()),
('Lasso', Lasso()),
('ElasticNet', ElasticNet()),
('Huber', HuberRegressor()),
('DT', DecisionTreeRegressor()),
('RF', RandomForestRegressor()),
('KNN', KNeighborsRegressor()),
('KernelRidge', KernelRidge()),
('MLP', MLPRegressor()),
('SVR', SVR())]
```

```
In [139]: num_folds = 5
seed = 7
```

```
In [143]: names = []
results = []

kfold = KFold(n_splits = num_folds, shuffle = True, random_state=seed)

for name, model in models: # 에러나면 .values.ravel() 빼고 y_train으로 해보기
    score = cross_val_score(model, X_train_scale, y_train.values.ravel(), cv =
kfold)
    names.append(name)
    results.append(score)
    print(name, score.mean().round(5))
```

```
Ridge 0.88372
Lasso 0.16192
ElasticNet 0.56008
Huber 0.87945
DT 0.9986
RF 0.99925
KNN 0.95446
KernelRidge -0.4884
MLP 0.98492
SVR 0.9748
```


```
In [142]: fig = plt.figure()
fig.suptitle('Algorithm Comparison')
ax = fig.add_subplot(111)
plt.boxplot(results)
ax.set_xticklabels(names)
plt.show()
```

C:\Users\50008313\AppData\Local\Continuum\anaconda3\lib\site-packages\matplotlib\backends\backend\_agg.py:238: RuntimeWarning: Glyph 8722 missing from current font.

font.set\_text(s, 0.0, flags=flags)

C:\Users\50008313\AppData\Local\Continuum\anaconda3\lib\site-packages\matplotlib\backends\backend\_agg.py:201: RuntimeWarning: Glyph 8722 missing from current font.

font.set\_text(s, 0, flags=flags)



===== 여기는 GridSearchCV 참고 내용

```
In [ ]: models = []  
        params = []
```

```
In [ ]: model = ('Ridge', Ridge())  
        param = {  
            'alpha': [0.1, 0.3, 0.5, 1.0, 3.0, 5.0, 10.0]  
        }  
  
        models.append(model)  
        params.append(param)
```

```
In [ ]: model = ('Lasso', Lasso())  
        param = {  
            'alpha': [0.1, 0.3, 0.5, 1.0, 3.0, 5.0, 10.0]  
        }  
  
        models.append(model)  
        params.append(param)
```

```
In [ ]: model = ('ElasticNet', ElasticNet())  
        param = {  
            'alpha': [0.1, 0.3, 0.5, 1.0, 3.0, 5.0, 10.0],  
            'l1_ratio': [0.3, 0.5, 0.7]  
        }  
  
        models.append(model)  
        params.append(param)
```

```
In [ ]: model = ('HuberReg', HuberRegressor())  
        param = {  
            'alpha': [0.0001, 0.001, 0.01]  
        }  
  
        models.append(model)  
        params.append(param)
```

```
In [ ]: model = ('CART', DecisionTreeRegressor())  
        param = {  
            'max_depth': [2, 3, 4, 5],  
            'min_samples_split': [0.02, 0.05]  
        }  
  
        models.append(model)  
        params.append(param)
```

```
In [ ]: model = ('RandomForest', RandomForestRegressor())
        param = {
            'n_estimators': [50, 60, 70, 80, 90, 100],
            'max_features': [6, 7, 8, 9, 10]
        }

        models.append(model)
        params.append(param)
```

```
In [ ]: model = ('KNN', KNeighborsRegressor())
        param = {
            'KNN__n_neighbors': [5, 10, 15, 20, 25, 30],
            'KNN__weights': ['uniform', 'distance']
        }

        models.append(model)
        params.append(param)
```

```
In [ ]: model = ('KernelRidge', KernelRidge())
        param = [
            {'kernel': ['linear'], 'alpha': [0.01, 0.05, 0.1, 0.5, 1.0]},
            {'kernel': ['rbf'], 'alpha': [0.01, 0.05, 0.1, 0.5, 1.0], 'gamma': [0.01,
            0.05, 0.1, 0.5, 1.0, 5.0, 10.0]}
        ]

        models.append(model)
        params.append(param)
```

```
In [ ]: model = ('MLP', MLPRegressor())
        param = {
            'hidden_layer_sizes': [(50, ), (100, ), (50, 50), (100, 100)],
            'solver': ['lbfgs'],
            'alpha': [0.0001, 0.001, 0.005],
            'max_iter': [200, 300, 400]
        }

        models.append(model)
        params.append(param)
```

```
In [ ]: model = ('SVR', SVR())
        param = [
            {'kernel': ['linear'], 'C': [1.0, 10.0, 50.0, 100.0]},
            {'kernel': ['rbf'], 'C': [1.0, 10.0, 50.0, 100.0], 'gamma': [0.01, 0.05,
            0.1, 0.5, 1.0]},
            {'kernel': ['poly'], 'C': [1.0, 10.0, 50.0, 100.0], 'degree': [3, 4, 5]}
        ]

        models.append(model)
        params.append(param)
```

## 파라미터 튜닝 및 최종 모델 선정

```
In [ ]: model = RandomForestRegressor()

n_estimators_set = [50, 60, 70, 80, 90, 100]
max_features_set = [6, 7, 8, 9, 10]
param_grid = dict(n_estimators = n_estimators_set,
                  max_features = max_features_set)

grid = GridSearchCV(estimator=model, param_grid=param_grid, cv=kfold)
grid_result = grid.fit(X_train_scale, y_train.values.ravel()) # 에러나면 .values.ravel() 빼고 y_train으로 해보기
print('Best : %f using %s' % (grid_result.best_score_, grid_result.best_params_))

a = grid_result.cv_results_

for i in range(len(a['rank_test_score'])):
    print('%f (%f) with: %r' % (a['mean_test_score'][i], a['std_test_score'][i], a['params'][i]))

# for params, mean_score, scores in grid_result.cv_results_: ## 애 에러난다
#     print('%f (%f) with: %r' % (mean_test_score.mean(), std_test_score.mean(), params))
```

```
In [156]: fine_tuned_RF = grid_result.best_estimator_
print('best params: ', grid_result.best_params_)
fine_tuned_RF.feature_importances_
```

```
best params: {'max_features': 7, 'n_estimators': 50}
```

```
Out[156]: array([0.00000000e+00, 3.16966872e-02, 1.46239923e-04, 1.14957090e-04,
1.11716386e-01, 8.56316024e-01, 9.70581592e-06])
```

```
In [157]: pd.DataFrame({'col':X_train_scale.columns, 'FI':fine_tuned_RF.feature_importances_}).sort_values('FI', ascending=False)
```


```
Out[157]:
```

	col	FI
5	Monetary	0.856316
4	Frequency	0.111716
1	temp	0.031697
2	humidity	0.000146
3	windspeed	0.000115
6	kmeans	0.000010
0	Intercept	0.000000



```
In [158]: importances = pd.Series(fine_tuned_RF.feature_importances_, index=X_train_scale.columns)
importances.plot(kind='bar')
```

Out[158]: <AxesSubplot:>



## Test set 활용하여 예측 수행

```
In [159]: y_pred = fine_tuned_RF.predict(X_test_scale)
```

```
In [160]: from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

## R<sup>2</sup>

```
In [161]: r2 = r2_score(np.exp1(y_test), np.exp1(y_pred)) # log처리 안 했으면 그냥 y_test, y_pred
r2
```

Out[161]: 0.9995663230063383

## MSE

```
In [162]: mse = mean_squared_error(np.exp1(y_test), np.exp1(y_pred)) # log처리 안 했으면 그냥 y_test, y_pred
mse
```

Out[162]: 0.0008766132338490894

## RMSE

```
In [163]: rmse = np.sqrt(mse)
          rmse
```

```
Out[163]: 0.029607654987335443
```

## MAE

```
In [164]: mae = mean_absolute_error(np.exp1(y_test), np.exp1(y_pred)) # Log처리 안 했으
          면 그냥 y_test, y_pred
          mae
```

```
Out[164]: 0.013092585562396663
```

## MAPE

```
In [165]: def mp(y_test, y_pred):
          y_test, y_pred = np.array(y_test), np.array(y_pred)
          return np.mean(np.abs(y_test - y_pred)/y_test) * 100
          # 평균 절대 백분율 오차(MAPE)는 정확도를 오차의 백분율로 표시합니다.
          # MAPE는 백분율이기 때문에 다른 정확도 측도 통계량보다 더 쉽게 이해할 수 있습니다.
          # 예를 들어 MAPE가 50이면 예측 값은 평균 5% 벗어납니다
```

```
In [166]: mape = mp(np.exp1(y_test), np.exp1(y_pred)) # Log처리 안 했으면 그냥 y_test, y_
          pred
          mape
```

```
Out[166]: 50.28324823737699
```


```
=====
```

```
In [197]: # 데이터를 일부러 줄임(2011년 8월 자료로만)
          data = dfts[(dfts.index.year == 2011)&(dfts.index.month == 8)]
```

```
In [198]: data.shape
```

```
Out[198]: (456, 1)
```

```
In [199]: fig = data.plot()
```



```
In [200]: # Seasonal decomposition plot : Seasonal decomposition using moving averages.


# Observed : observed data
# Trend : The estimated trend component
# Seasonal : The estimated seasonal component
# resid : The estimated residuals
decomposition = sm.tsa.seasonal_decompose(data['count'], model = 'additive', p
period=1)
fig = decomposition.plot()
fig.set_size_inches(10,10)
plt.show()
```

C:\Users\50008313\AppData\Local\Continuum\anaconda3\lib\site-packages\matplotlib\backends\backend\_agg.py:238: RuntimeWarning: Glyph 8722 missing from current font.

font.set\_text(s, 0.0, flags=flags)

C:\Users\50008313\AppData\Local\Continuum\anaconda3\lib\site-packages\matplotlib\backends\backend\_agg.py:201: RuntimeWarning: Glyph 8722 missing from current font.

font.set\_text(s, 0, flags=flags)



## Train. test set split

```
In [201]: # Tr, Te = 8:2
train_data, test_data = train_test_split(data, test_size=0.2, shuffle=False)
```

## 정상성 확인

```
In [203]: # ACF, PACF plot

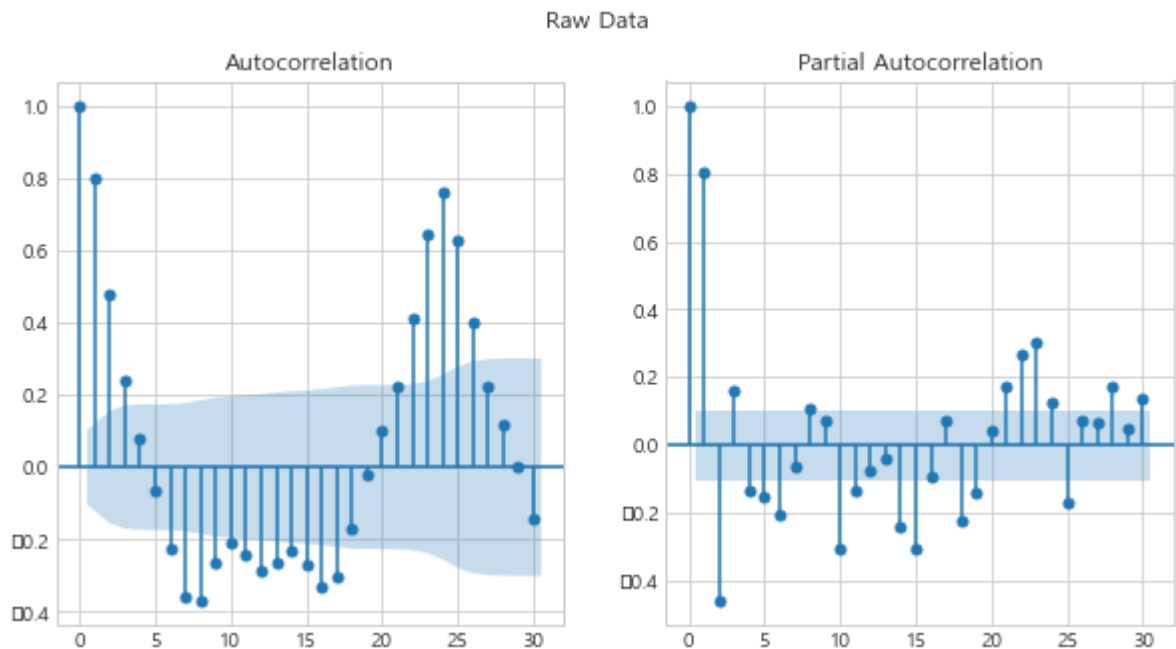
fig, ax = plt.subplots(1,2, figsize = (10, 5))
fig.suptitle('Raw Data')
sm.graphics.tsa.plot_acf(train_data.values.squeeze(), lags = 30, ax = ax[0])
sm.graphics.tsa.plot_pacf(train_data.values.squeeze(), lags = 30, ax = ax[1])
plt.show()
### ACF 그래프가 점진적으로 감소하는 것은 전형적인 Non-stationary 데이터이다 = 정상성이 없음
```

C:\Users\50008313\AppData\Local\Continuum\anaconda3\lib\site-packages\matplotlib\lib\backends\backend\_agg.py:238: RuntimeWarning: Glyph 8722 missing from current font.

font.set\_text(s, 0.0, flags=flags)

C:\Users\50008313\AppData\Local\Continuum\anaconda3\lib\site-packages\matplotlib\lib\backends\backend\_agg.py:201: RuntimeWarning: Glyph 8722 missing from current font.

font.set\_text(s, 0, flags=flags)



## 차분

In [205]: *# Differencing*

```

diff_train_data = train_data.copy()
diff_train_data = diff_train_data['count'].diff()
diff_train_data = diff_train_data.dropna()
print('##### Raw Data #####')
print(train_data)
print('### Differenced Data ###')
print(diff_train_data)

```

##### Raw Data #####

	count
date	
2011-08-01 00:00:00	29.0
2011-08-01 01:00:00	17.0
2011-08-01 02:00:00	11.0
2011-08-01 03:00:00	4.0
2011-08-01 04:00:00	4.0
...	...
2011-08-15 23:00:00	88.0
2011-08-16 00:00:00	31.0
2011-08-16 01:00:00	16.0
2011-08-16 02:00:00	4.0
2011-08-16 03:00:00	6.0

[364 rows x 1 columns]

### Differenced Data ###

date	
2011-08-01 01:00:00	-12.0
2011-08-01 02:00:00	-6.0
2011-08-01 03:00:00	-7.0
2011-08-01 04:00:00	0.0
2011-08-01 05:00:00	22.0
...	...
2011-08-15 23:00:00	-10.0
2011-08-16 00:00:00	-57.0
2011-08-16 01:00:00	-15.0
2011-08-16 02:00:00	-12.0
2011-08-16 03:00:00	2.0

Name: count, Length: 363, dtype: float64

In [207]: *# differenced data plot*


```
plt.figure(figsize = (12,8))
plt.subplot(211)
plt.plot(train_data['count'])
plt.subplot(212)
plt.plot(diff_train_data, 'orange') # first difference (t - (t-1))
plt.legend(['Differenced Data (Stationary)'])
plt.show()
```

C:\Users\50008313\AppData\Local\Continuum\anaconda3\lib\site-packages\matplotlib\backends\backend\_agg.py:238: RuntimeWarning: Glyph 8722 missing from current font.

font.set\_text(s, 0.0, flags=flags)

C:\Users\50008313\AppData\Local\Continuum\anaconda3\lib\site-packages\matplotlib\backends\backend\_agg.py:201: RuntimeWarning: Glyph 8722 missing from current font.

font.set\_text(s, 0, flags=flags)



In [208]: *# ACF, PACF plot*

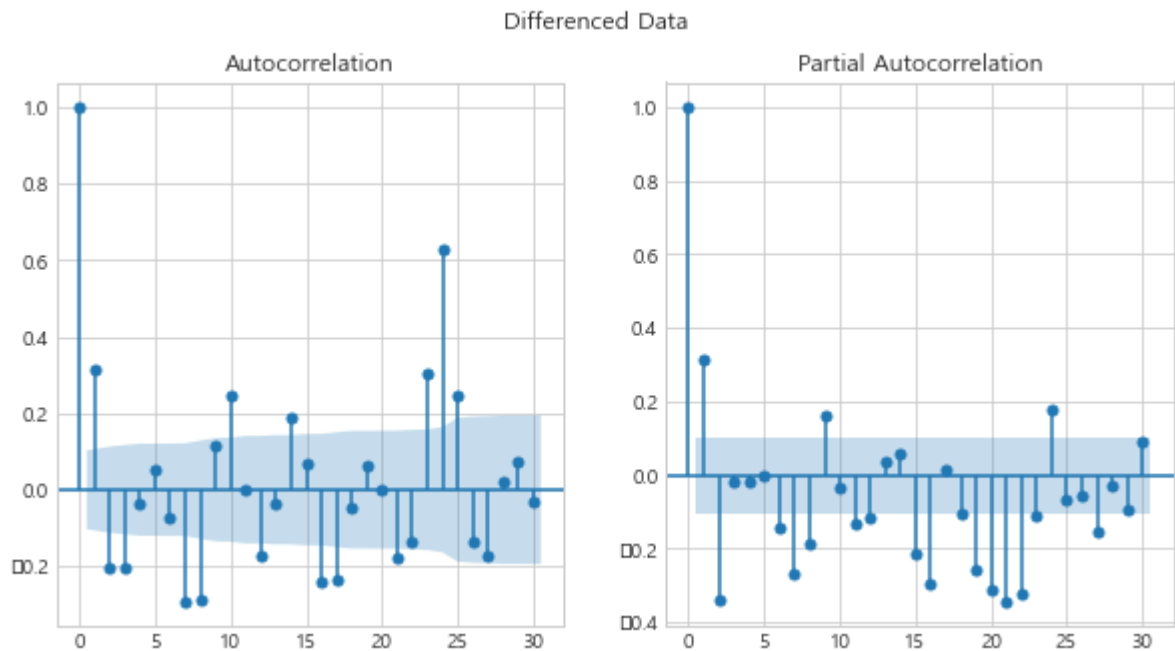
```
fig, ax = plt.subplots(1, 2, figsize = (10, 5))
fig.suptitle('Differenced Data')
sm.graphics.tsa.plot_acf(diff_train_data.values.squeeze(), lags = 30, ax = ax[0])
sm.graphics.tsa.plot_pacf(diff_train_data.values.squeeze(), lags = 30, ax = ax[1])
plt.show()
```

C:\Users\50008313\AppData\Local\Continuum\anaconda3\lib\site-packages\matplotlib\backends\backend\_agg.py:238: RuntimeWarning: Glyph 8722 missing from current font.

font.set\_text(s, 0.0, flags=flags)

C:\Users\50008313\AppData\Local\Continuum\anaconda3\lib\site-packages\matplotlib\backends\backend\_agg.py:201: RuntimeWarning: Glyph 8722 missing from current font.

font.set\_text(s, 0, flags=flags)



## 기본 모델 생성



```
In [209]: # ARIMA model fitting
# The (p, d, q) order of the model for the number of AR parameters, difference
s, and MA parameters to use.

model = ARIMA(train_data.values, order=(1,1,1))
model_fit = model.fit()
model_fit.summary()

# AIC 값은 1069.440이고, constant의 p-value 값이 유의미하지 않게 나왔다.
```

Out[209]:

ARIMA Model Results

<b>Dep. Variable:</b>	D.y	<b>No. Observations:</b>	363
<b>Model:</b>	ARIMA(1, 1, 1)	<b>Log Likelihood</b>	-2122.751
<b>Method:</b>	css-mle	<b>S.D. of innovations</b>	83.807
<b>Date:</b>	Wed, 09 Dec 2020	<b>AIC</b>	4253.502
<b>Time:</b>	22:54:27	<b>BIC</b>	4269.079
<b>Sample:</b>	1	<b>HQIC</b>	4259.694

	coef	std err	z	P> z	[0.025	0.975]
<b>const</b>	-0.0571	6.142	-0.009	0.993	-12.095	11.981
<b>ar.L1.D.y</b>	-0.1120	0.102	-1.094	0.274	-0.313	0.089
<b>ma.L1.D.y</b>	0.5537	0.084	6.606	0.000	0.389	0.718

Roots

	Real	Imaginary	Modulus	Frequency
<b>AR.1</b>	-8.9287	+0.0000j	8.9287	0.5000
<b>MA.1</b>	-1.8059	+0.0000j	1.8059	0.5000

```
In [210]: model = ARIMA(train_data.values, order=(0,1,1))
          model_fit = model.fit()
          model_fit.summary()
```

Out[210]: ARIMA Model Results

<b>Dep. Variable:</b>	D.y	<b>No. Observations:</b>	363
<b>Model:</b>	ARIMA(0, 1, 1)	<b>Log Likelihood</b>	-2123.378
<b>Method:</b>	css-mle	<b>S.D. of innovations</b>	83.952
<b>Date:</b>	Wed, 09 Dec 2020	<b>AIC</b>	4252.756
<b>Time:</b>	23:04:33	<b>BIC</b>	4264.439
<b>Sample:</b>	1	<b>HQIC</b>	4257.400

	coef	std err	z	P> z	[0.025	0.975]
<b>const</b>	-0.0619	6.482	-0.010	0.992	-12.766	12.642
<b>ma.L1.D.y</b>	0.4724	0.047	10.044	0.000	0.380	0.565

Roots

	Real	Imaginary	Modulus	Frequency
<b>MA.1</b>	-2.1170	+0.0000j	2.1170	0.5000

```
In [211]: model = ARIMA(train_data.values, order=(1,1,0))
          model_fit = model.fit()
          model_fit.summary()
```

Out[211]: ARIMA Model Results

<b>Dep. Variable:</b>	D.y	<b>No. Observations:</b>	363
<b>Model:</b>	ARIMA(1, 1, 0)	<b>Log Likelihood</b>	-2136.885
<b>Method:</b>	css-mle	<b>S.D. of innovations</b>	87.152
<b>Date:</b>	Wed, 09 Dec 2020	<b>AIC</b>	4279.769
<b>Time:</b>	23:04:38	<b>BIC</b>	4291.453
<b>Sample:</b>	1	<b>HQIC</b>	4284.413

	coef	std err	z	P> z	[0.025	0.975]
<b>const</b>	-0.0760	6.640	-0.011	0.991	-13.089	12.937
<b>ar.L1.D.y</b>	0.3119	0.050	6.267	0.000	0.214	0.409

Roots

	Real	Imaginary	Modulus	Frequency
<b>AR.1</b>	3.2060	+0.0000j	3.2060	0.0000

## 최적 모델 탐색

```
In [213]: print('Examples of parameter combinations for Seasonal ARIMA...')
p = range(0,3)
d = range(1,2)
q = range(0,3)

pdq = list(itertools.product(p, d, q))
pdq
```

Examples of parameter combinations for Seasonal ARIMA...

```
Out[213]: [(0, 1, 0),
           (0, 1, 1),
           (0, 1, 2),
           (1, 1, 0),
           (1, 1, 1),
           (1, 1, 2),
           (2, 1, 0),
           (2, 1, 1),
           (2, 1, 2)]
```

```
In [214]: aic=[]
for i in pdq:
    model = ARIMA(train_data.values, order=(i))
    model_fit = model.fit()
    print(f'ARIMA: {i} >> AIC : {round(model_fit.aic, 2)}')
    aic.append(round(model_fit.aic,2))
```

```
ARIMA: (0, 1, 0) >> AIC : 4315.04
ARIMA: (0, 1, 1) >> AIC : 4252.76
ARIMA: (0, 1, 2) >> AIC : 4252.07
ARIMA: (1, 1, 0) >> AIC : 4279.77
ARIMA: (1, 1, 1) >> AIC : 4253.5
ARIMA: (1, 1, 2) >> AIC : 4198.08
ARIMA: (2, 1, 0) >> AIC : 4238.56
ARIMA: (2, 1, 1) >> AIC : 4240.4
ARIMA: (2, 1, 2) >> AIC : 4194.67
```

```
In [215]: # Search optimal parameters

optimal = [(pdq[i], j) for i, j in enumerate(aic) if j == min(aic)]
optimal
```

```
Out[215]: [(2, 1, 2), 4194.67]
```

In [218]: # 위 최적 값으로 만든 모델 다시 Summary

```
model_opt = ARIMA(train_data.values, order = optimal[0][0])
model_opt_fit = model_opt.fit()
model_opt_fit.summary()
```

# AIC score가 1045.66으로 임의의 모델보다 성능이 좋아졌고, p-value도 모두 유의미하게 나옴

Out[218]:

ARIMA Model Results

<b>Dep. Variable:</b>	D.y	<b>No. Observations:</b>	363
<b>Model:</b>	ARIMA(2, 1, 2)	<b>Log Likelihood</b>	-2091.334
<b>Method:</b>	css-mle	<b>S.D. of innovations</b>	76.390
<b>Date:</b>	Wed, 09 Dec 2020	<b>AIC</b>	4194.669
<b>Time:</b>	23:12:34	<b>BIC</b>	4218.035
<b>Sample:</b>	1	<b>HQIC</b>	4203.957

	coef	std err	z	P> z	[0.025	0.975]
<b>const</b>	0.0155	0.154	0.101	0.920	-0.286	0.317
<b>ar.L1.D.y</b>	0.8791	0.097	9.048	0.000	0.689	1.070
<b>ar.L2.D.y</b>	-0.2176	0.088	-2.465	0.014	-0.391	-0.045
<b>ma.L1.D.y</b>	-0.6179	0.094	-6.608	0.000	-0.801	-0.435
<b>ma.L2.D.y</b>	-0.3821	0.093	-4.096	0.000	-0.565	-0.199

Roots

	Real	Imaginary	Modulus	Frequency
<b>AR.1</b>	2.0203	-0.7174j	2.1439	-0.0543
<b>AR.2</b>	2.0203	+0.7174j	2.1439	0.0543
<b>MA.1</b>	1.0000	+0.0000j	1.0000	0.0000
<b>MA.2</b>	-2.6172	+0.0000j	2.6172	0.5000

## Test 데이터 예측

```
In [219]: prediction = model_opt_fit.forecast(len(test_data))
predicted_value = prediction[0] # predicted_value 가 y_pred
predicted_ub = prediction[2][:,0]
predicted_lb = prediction[2][:,1]
predict_index = list(test_data.index)
r2 = r2_score(test_data, predicted_value)
```

```
In [222]: fig, ax = plt.subplots(figsize=(12,6))

ax.plot(predict_index, predicted_value, color = 'orange', label = 'Prediction'
) # 예측값(위 vline 이후 구간에 표시됨)
ax.fill_between(predict_index, predicted_lb, predicted_ub, color = 'k', alpha
= 0.1, label = '0.95 Prediction Interval')

data.plot(ax = ax);
# ax.vlines('1958-08-01', 0, 1000, linestyle = '--', color = 'r', label = 'Sta
rt of Forecast') # x좌표를 날짜로 적음
# ax.legend(loc='upper left')
# plt.suptitle(f'ARIMA {optimal[0][0]} Prediction Results (r2_score: {round(r
2,2)})')

plt.show()


# 빨간 점선 이후의 주황색 선이 예측값이며, 회색 구간이 95% interval 구간이다.
# 대체로 추세를 따라가나 피크 값을 완벽히 예측하기에는 다소 무리가 있는 것을 볼 수 있다.
# R2 score도 0.22 수준인 것을 확인할 수 있었다.
```

C:\Users\50008313\AppData\Local\Continuum\anaconda3\lib\site-packages\matplotlib\backends\backend\_agg.py:238: RuntimeWarning: Glyph 8722 missing from current font.

font.set\_text(s, 0.0, flags=flags)

C:\Users\50008313\AppData\Local\Continuum\anaconda3\lib\site-packages\matplotlib\backends\backend\_agg.py:201: RuntimeWarning: Glyph 8722 missing from current font.

font.set\_text(s, 0, flags=flags)



## 성능 평가

```
In [160]: from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

## R<sup>2</sup>

```
In [223]: r2 = r2_score(test_data, predicted_value)
r2
```

Out[223]: 0.019727875700680686

## MSE

```
In [224]: mse = mean_squared_error(test_data, predicted_value)
mse
```

Out[224]: 20910.821211736336

## RMSE

```
In [225]: rmse = np.sqrt(mse)
rmse
```

Out[225]: 144.60574404820971

## MAE

```
In [226]: mae = mean_absolute_error(test_data, predicted_value)
mae
```

Out[226]: 108.6861962191569

## MAPE

```
In [165]: def mp(y_test, y_pred):
    y_test, y_pred = np.array(y_test), np.array(y_pred)
    return np.mean(np.abs(y_test - y_pred)/y_test) * 100
    # 평균 절대 백분율 오차(MAPE)는 정확도를 오차의 백분율로 표시합니다.
    # MAPE는 백분율이기 때문에 다른 정확도 측도 통계량보다 더 쉽게 이해할 수 있습니다.
    # 예를 들어 MAPE가 5이면 예측 값은 평균 5% 벗어납니다
```

```
In [227]: mape = mp(test_data, predicted_value)
mape
```

Out[227]: 441.99048769365373

## 데이터 생성

```
In [244]: # len(test_data) 자리에 원하는 만큼 숫자 넣어주면 됨

prediction = model_opt_fit.forecast(len(test_data))
predicted_value = prediction[0] # predicted_value 가 y_pred
predicted_ub = prediction[2][:,0]
predicted_lb = prediction[2][:,1]
predict_index = list(test_data.index)
r2 = r2_score(test_data, predicted_value)
```

Out[244]: Timestamp('2011-08-19 23:00:00')

## (참고) 계절성 반영 시 SARIMA 모델링 수행

```
In [230]: print('Examples of parameter combinations for Seasonal ARIMA...')
p = range(0,3)
d = range(1,2)
q = range(0,3)
pdq = list(itertools.product(p, d, q))
seasonal_pdq = [(x[0], x[1], x[2], 7 # 계절성이 7일마다 있다고 생각해서 7 입력. 애도
위에 range로 탐색해도 됨
) for x in list(itertools.product(p, d, q))]
seasonal_pdq
```

Examples of parameter combinations for Seasonal ARIMA...

Out[230]: [(0, 1, 0, 7),  
(0, 1, 1, 7),  
(0, 1, 2, 7),  
(1, 1, 0, 7),  
(1, 1, 1, 7),  
(1, 1, 2, 7),  
(2, 1, 0, 7),  
(2, 1, 1, 7),  
(2, 1, 2, 7)]

```
In [231]: aic = []
          params = []
          for i in pdq:
              for j in seasonal_pdq:
                  try:
                      model = SARIMAX(train_data.values, order=(i), seasonal_order = (j
                      ))
                      model_fit = model.fit()
                      print(f'SARIMA: {i}{j} >> AIC : {round(model_fit.aic,2)}')
                      aic.append(round(model_fit.aic,2))
                      params.append((i, j))
                  except:
                      continue
```



```

SARIMA: (0, 1, 0)(0, 1, 0, 7) >> AIC : 4572.25
SARIMA: (0, 1, 0)(0, 1, 1, 7) >> AIC : 4261.5
SARIMA: (0, 1, 0)(0, 1, 2, 7) >> AIC : 4236.78
SARIMA: (0, 1, 0)(1, 1, 0, 7) >> AIC : 4346.58
SARIMA: (0, 1, 0)(1, 1, 1, 7) >> AIC : 4230.88
SARIMA: (0, 1, 0)(1, 1, 2, 7) >> AIC : 4227.12
SARIMA: (0, 1, 0)(2, 1, 0, 7) >> AIC : 4324.8
SARIMA: (0, 1, 0)(2, 1, 1, 7) >> AIC : 4228.58
SARIMA: (0, 1, 0)(2, 1, 2, 7) >> AIC : 4228.98
SARIMA: (0, 1, 1)(0, 1, 0, 7) >> AIC : 4472.44
SARIMA: (0, 1, 1)(0, 1, 1, 7) >> AIC : 4197.85
SARIMA: (0, 1, 1)(0, 1, 2, 7) >> AIC : 4190.14
SARIMA: (0, 1, 1)(1, 1, 0, 7) >> AIC : 4286.48
SARIMA: (0, 1, 1)(1, 1, 1, 7) >> AIC : 4187.24
SARIMA: (0, 1, 1)(1, 1, 2, 7) >> AIC : 4196.09
SARIMA: (0, 1, 1)(2, 1, 0, 7) >> AIC : 4268.11
SARIMA: (0, 1, 1)(2, 1, 1, 7) >> AIC : 4182.24

```

```

C:\Users\50008313\AppData\Local\Continuum\anaconda3\lib\site-packages\statsmodels\base\model.py:568: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals

```

```

    "Check mle_retvals", ConvergenceWarning)

```

```

SARIMA: (0, 1, 1)(2, 1, 2, 7) >> AIC : 4189.68
SARIMA: (0, 1, 2)(0, 1, 0, 7) >> AIC : 4474.17
SARIMA: (0, 1, 2)(0, 1, 1, 7) >> AIC : 4197.96
SARIMA: (0, 1, 2)(0, 1, 2, 7) >> AIC : 4188.07
SARIMA: (0, 1, 2)(1, 1, 0, 7) >> AIC : 4288.25
SARIMA: (0, 1, 2)(1, 1, 1, 7) >> AIC : 4185.4
SARIMA: (0, 1, 2)(1, 1, 2, 7) >> AIC : 4194.15
SARIMA: (0, 1, 2)(2, 1, 0, 7) >> AIC : 4270.03
SARIMA: (0, 1, 2)(2, 1, 1, 7) >> AIC : 4182.81
SARIMA: (0, 1, 2)(2, 1, 2, 7) >> AIC : 4186.93
SARIMA: (1, 1, 0)(0, 1, 0, 7) >> AIC : 4519.98
SARIMA: (1, 1, 0)(0, 1, 1, 7) >> AIC : 4224.81
SARIMA: (1, 1, 0)(0, 1, 2, 7) >> AIC : 4210.56
SARIMA: (1, 1, 0)(1, 1, 0, 7) >> AIC : 4297.35
SARIMA: (1, 1, 0)(1, 1, 1, 7) >> AIC : 4205.15
SARIMA: (1, 1, 0)(1, 1, 2, 7) >> AIC : 4220.77
SARIMA: (1, 1, 0)(2, 1, 0, 7) >> AIC : 4282.27
SARIMA: (1, 1, 0)(2, 1, 1, 7) >> AIC : 4197.44

```

```

C:\Users\50008313\AppData\Local\Continuum\anaconda3\lib\site-packages\statsmodels\base\model.py:568: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals

```

```

    "Check mle_retvals", ConvergenceWarning)

```

```

SARIMA: (1, 1, 0)(2, 1, 2, 7) >> AIC : 4207.81
SARIMA: (1, 1, 1)(0, 1, 0, 7) >> AIC : 4474.34
SARIMA: (1, 1, 1)(0, 1, 1, 7) >> AIC : 4198.97
SARIMA: (1, 1, 1)(0, 1, 2, 7) >> AIC : 4190.22
SARIMA: (1, 1, 1)(1, 1, 0, 7) >> AIC : 4288.31
SARIMA: (1, 1, 1)(1, 1, 1, 7) >> AIC : 4187.32
SARIMA: (1, 1, 1)(1, 1, 2, 7) >> AIC : 4196.35
SARIMA: (1, 1, 1)(2, 1, 0, 7) >> AIC : 4270.07
SARIMA: (1, 1, 1)(2, 1, 1, 7) >> AIC : 4183.45
SARIMA: (1, 1, 1)(2, 1, 2, 7) >> AIC : 4189.44
SARIMA: (1, 1, 2)(0, 1, 0, 7) >> AIC : 4421.86
SARIMA: (1, 1, 2)(0, 1, 1, 7) >> AIC : 4150.66
SARIMA: (1, 1, 2)(0, 1, 2, 7) >> AIC : 4141.32
SARIMA: (1, 1, 2)(1, 1, 0, 7) >> AIC : 4287.23
SARIMA: (1, 1, 2)(1, 1, 1, 7) >> AIC : 4139.87
SARIMA: (1, 1, 2)(1, 1, 2, 7) >> AIC : 4162.61
SARIMA: (1, 1, 2)(2, 1, 0, 7) >> AIC : 4226.17
SARIMA: (1, 1, 2)(2, 1, 1, 7) >> AIC : 4140.72

```

```

C:\Users\50008313\AppData\Local\Continuum\anaconda3\lib\site-packages\statsmodels\base\model.py:568: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals

```

```

    "Check mle_retvals", ConvergenceWarning)

```

```

SARIMA: (1, 1, 2)(2, 1, 2, 7) >> AIC : 4142.36
SARIMA: (2, 1, 0)(0, 1, 0, 7) >> AIC : 4449.53
SARIMA: (2, 1, 0)(0, 1, 1, 7) >> AIC : 4183.9
SARIMA: (2, 1, 0)(0, 1, 2, 7) >> AIC : 4177.57
SARIMA: (2, 1, 0)(1, 1, 0, 7) >> AIC : 4284.58
SARIMA: (2, 1, 0)(1, 1, 1, 7) >> AIC : 4175.94
SARIMA: (2, 1, 0)(1, 1, 2, 7) >> AIC : 4182.55
SARIMA: (2, 1, 0)(2, 1, 0, 7) >> AIC : 4256.46
SARIMA: (2, 1, 0)(2, 1, 1, 7) >> AIC : 4175.27
SARIMA: (2, 1, 0)(2, 1, 2, 7) >> AIC : 4179.2
SARIMA: (2, 1, 1)(0, 1, 0, 7) >> AIC : 4450.49
SARIMA: (2, 1, 1)(0, 1, 1, 7) >> AIC : 4185.62
SARIMA: (2, 1, 1)(0, 1, 2, 7) >> AIC : 4145.43
SARIMA: (2, 1, 1)(1, 1, 0, 7) >> AIC : 4245.92
SARIMA: (2, 1, 1)(1, 1, 1, 7) >> AIC : 4142.92
SARIMA: (2, 1, 1)(1, 1, 2, 7) >> AIC : 4163.12

```

```

C:\Users\50008313\AppData\Local\Continuum\anaconda3\lib\site-packages\statsmodels\base\model.py:568: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals

```

```

    "Check mle_retvals", ConvergenceWarning)

```

```

SARIMA: (2, 1, 1)(2, 1, 0, 7) >> AIC : 4220.79
SARIMA: (2, 1, 1)(2, 1, 1, 7) >> AIC : 4142.27
SARIMA: (2, 1, 1)(2, 1, 2, 7) >> AIC : 4142.98
SARIMA: (2, 1, 2)(0, 1, 0, 7) >> AIC : 4388.5
SARIMA: (2, 1, 2)(0, 1, 1, 7) >> AIC : 4181.84
SARIMA: (2, 1, 2)(0, 1, 2, 7) >> AIC : 4170.11

```

```

C:\Users\50008313\AppData\Local\Continuum\anaconda3\lib\site-packages\statsmodels\base\model.py:568: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals

```

```

    "Check mle_retvals", ConvergenceWarning)

```

```
SARIMA: (2, 1, 2)(1, 1, 0, 7) >> AIC : 4246.16
SARIMA: (2, 1, 2)(1, 1, 1, 7) >> AIC : 4165.27
SARIMA: (2, 1, 2)(1, 1, 2, 7) >> AIC : 4152.32
SARIMA: (2, 1, 2)(2, 1, 0, 7) >> AIC : 4253.74
SARIMA: (2, 1, 2)(2, 1, 1, 7) >> AIC : 4153.63
SARIMA: (2, 1, 2)(2, 1, 2, 7) >> AIC : 4139.82
```

In [232]: *# Search optimal parameters*

```
optimal = [(params[i], j) for i, j in enumerate(aic) if j == min(aic)]
optimal
```

*# small pdq는 (1,1,0), large pdq는 (1,1,2) 그리고 Seasonal parameter는 12인 것을 볼 수 있다.*

Out[232]: [(((2, 1, 2), (2, 1, 2, 7)), 4139.82)]

```
In [233]: model_opt = SARIMAX(train_data.values, order=optimal[0][0][0], seasonal_order
= optimal[0][0][1])
model_opt_fit = model_opt.fit()
model_opt_fit.summary()

# ARIMA보다 SARIMA가 AIC가 훨씬 낮은 것을 볼 수 있다.
```

Out[233]:

SARIMAX Results

<b>Dep. Variable:</b>	y	<b>No. Observations:</b>	364
<b>Model:</b>	SARIMAX(2, 1, 2)x(2, 1, 2, 7)	<b>Log Likelihood</b>	-2060.909
<b>Date:</b>	Wed, 09 Dec 2020	<b>AIC</b>	4139.819
<b>Time:</b>	23:24:07	<b>BIC</b>	4174.693
<b>Sample:</b>	0	<b>HQIC</b>	4153.691
	- 364		
<b>Covariance Type:</b>	opg		

	coef	std err	z	P> z	[0.025	0.975]
<b>ar.L1</b>	0.9201	0.135	6.834	0.000	0.656	1.184
<b>ar.L2</b>	-0.2209	0.131	-1.682	0.092	-0.478	0.036
<b>ma.L1</b>	-0.6882	2.654	-0.259	0.795	-5.889	4.513
<b>ma.L2</b>	-0.3115	0.873	-0.357	0.721	-2.023	1.400
<b>ar.S.L7</b>	-0.6366	0.572	-1.112	0.266	-1.759	0.485
<b>ar.S.L14</b>	-0.0065	0.146	-0.045	0.964	-0.293	0.279
<b>ma.S.L7</b>	-0.5230	0.777	-0.674	0.501	-2.045	0.999
<b>ma.S.L14</b>	-0.4729	0.587	-0.806	0.420	-1.623	0.677
<b>sigma2</b>	5610.3109	1.45e+04	0.386	0.700	-2.29e+04	3.41e+04

<b>Ljung-Box (Q):</b>	235.96	<b>Jarque-Bera (JB):</b>	124.42
<b>Prob(Q):</b>	0.00	<b>Prob(JB):</b>	0.00
<b>Heteroskedasticity (H):</b>	0.77	<b>Skew:</b>	0.73
<b>Prob(H) (two-sided):</b>	0.16	<b>Kurtosis:</b>	5.50

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

## Test data 예측 및 성능평가는 위와 동일

```
In [234]: prediction = model_opt_fit.get_forecast(len(test_data))
predicted_value = prediction.predicted_mean
predicted_ub = prediction.conf_int()[0]
predicted_lb = prediction.conf_int()[1]
predict_index = list(test_data.index)
r2 = r2_score(test_data, predicted_value)
```

```
In [236]: fig, ax = plt.subplots(figsize=(12,6))

ax.plot(predict_index, predicted_value, color = 'orange', label = 'Prediction'
) # 예측값(위 vline 이후 구간에 표시됨)
ax.fill_between(predict_index, predicted_lb, predicted_ub, color = 'k', alpha
= 0.1, label = '0.95 Prediction Interval')

data.plot(ax = ax);
# ax.vlines('1958-08-01', 0, 700, linestyle = '--', color = 'r', label = 'Star
t of Forecast') # x좌표를 날짜로 적음
# ax.legend(loc='upper left')
# plt.suptitle(f'SARIMA {optimal[0][0][0]},{optimal[0][0][1]} Prediction Resul
ts (r2_score: {round(r2,2)})')

plt.show()


# 예측 값의 추세가 실제 값을 상당히 잘 따라가고 있으며,
# r2 score가 0.89로 훨씬 더 성능이 향상됨
# 계절성을 반영한 것이 예측 성능을 향상시키는데 기여를 했다고 볼 수 있다.
```

C:\Users\50008313\AppData\Local\Continuum\anaconda3\lib\site-packages\matplotlib\backends\backend\_agg.py:238: RuntimeWarning: Glyph 8722 missing from current font.

font.set\_text(s, 0.0, flags=flags)

C:\Users\50008313\AppData\Local\Continuum\anaconda3\lib\site-packages\matplotlib\backends\backend\_agg.py:201: RuntimeWarning: Glyph 8722 missing from current font.

font.set\_text(s, 0, flags=flags)



## Test data 예측 및 성능평가는 위와 동일